BlueDot Impact Community Lead Work Test

Curated Talent Marketplace for AI Safety

Eitan Sprejer / eitan@baish.com.ar

Links:

- Google Doc (live version): Full Document
- Loom walkthrough: Video
- Presentation slides: PDF on Google Drive
- Airtable base: BlueDot Talent Marketplace
- Repository: github.com/Eitan-Sprejer/bluedot-work-test

Executive Summary

The Problem

BlueDot has trained 7,000+ people. Some are exceptional. But there's no fast track to get the best of them into impactful roles, funded projects, or collaborations. Meanwhile, safety orgs can't find senior talent, funders have more money than good applications, and 10-30% of fellowship alumni cycle into another fellowship rather than converting to work.

The Solution

A monthly, invite-only event in SF (~50 people, ~3 hours) that puts BlueDot's top talent in a room with the institutional leaders who can hire them, fund them, or point them in the right direction. Ten structured 15-minute 1-on-1s per person, with pre-matching based on profiles. Think of it as a higher-signal, more exclusive EAG that happens every month.

Why It Works

We're not reinventing the wheel. EAG events already prove that curated 1-on-1 matching produces career outcomes (~10 meaningful connections per attendee, $70/connection). We take that proven mechanic and make it more exclusive (top 1% talent only), more frequent (monthly), and more accountable (dedicated facilitator follow-up with every participant).

What Makes BlueDot Uniquely Positioned

Two assets no one else has: (1) a curated talent database with performance data on 7,000+ alumni, and (2) direct relationships with the funders, safety orgs, and programs on the demand side. The combination is the moat.

The Specific Outcome

Every attendee leaves with at least one concrete opportunity they wouldn't have found otherwise. Measured by: job interviews, offers, and hires traced to event connections; funding applications submitted; collaborations started. Tracked through facilitator follow-up calls, not just surveys.

The System

An Airtable base (link) tracks events, next steps, feedback, outcomes, and individual participant progress. A five-layer system handles operations (Swapcard + Luma), data collection (forms), participant tracking (facilitators), outcome measurement, and monthly improvement. Every event generates data that makes the next one better.

My Approach

I spent the majority of my time on problem diagnosis: researching the AI safety talent pipeline to identify the specific bottleneck this event should address. I believe the Community Lead's most important job is knowing *why* an event exists, not just how to run one. The theory of impact in Part 1 is the core of this submission.

Part 1: Event Design

1. Theory of Impact

The Problem: No Fast Track for Top Talent

BlueDot knows where the best talent is. There's no fast track to get them doing impactful work.

BlueDot's pipeline has produced 7,000+ trained, motivated people. Some of them are exceptional: strong track records, deep skills, and ideas for projects, startups, or research that

could meaningfully reduce AI risk. BlueDot has performance data on these people (course evaluations, project quality, facilitator recommendations). They know who the top 1% are.

But knowing who the talent is and activating that talent are different things. Right now, a top-performing BlueDot alumnus finishes their course and... what? They go back to their existing job. They browse 80,000 Hours listings. They might apply to something. Many don't. The pipeline leaks precisely at the moment where the most promising people need a push.

This is not a hypothetical problem. The data shows it clearly:

1. **Safety orgs can't find the people they need.** After interviewing 10+ AI safety hiring managers, BlueDot's Li-Lian Ang found the core challenge is "sourcing talent who can own and lead work independently, and care deeply about AI risk." These people exist in BlueDot's database. The orgs just don't know about them. ("We asked 10+ AI safety orgs about their hiring needs", Feb 2026)

2. **Funders have more money than good applications.** Coefficient Giving (formerly Open Philanthropy) gave away >$150M in 2025 and explicitly says there's a "shortage of capacity-building interest relative to funding appetite." They want more people to apply. The bottleneck is not funding, it's people knowing what to apply for and having the confidence to do it. ("The case for AI safety capacity-building work", March 2026)

3. **Training doesn't automatically convert to impact.** Chris Clay analyzed 600+ alumni from 9 major AI safety fellowships and found that 10-30% cycle into another fellowship rather than converting to employment or projects. The pipeline produces trained people but lacks a bridge between training and doing. (EA Forum, 2026)

4. **The senior talent bottleneck compounds everything.** A March 2026 study (Teichman, 23 interviews with hiring managers and funders) found that "AI safety organizations are capacity constrained by a lack of senior researchers who can mentor and supervise junior talent." Without enough senior people, junior talent can't develop. And AI automation is raising the bar for what humans need to contribute, not lowering it. (EA Forum)

The talent exists. The opportunities exist. They aren't meeting efficiently. There is no dedicated, recurring, high-signal mechanism for connecting BlueDot's best people with the orgs, funders, and programs that need them.

Why This Works: Don't Reinvent the Wheel

We don't need to invent a new format. EAG events already prove that structured 1-on-1 matching drives real career outcomes:

- EAGx events achieve **~10 new connections per attendee** at a **cost per connection of ~$70** (EAGxAmsterdam 2025 retrospective, EA Forum)
- EAGxCDMX 2025: **1,714 total new connections** across 182 attendees
- Co-founder matches, job placements, and research collaborations trace directly back to EAG 1-on-1s
- 80,000 Hours' survey of 38 AI safety org leaders ranked **"attend events and get to know people"** as one of the top two pieces of career advice (Sept 2025)
- EA North 2025 (35 people, one day, GBP 50/person) achieved **4.56 new connections per person** and was rated **8.8/10** by attendees, making it **12x more cost-effective than flagship EAG** per connection

The evidence is clear: curated people in a room with structured matching produces career outcomes. The question is how to maximize the signal.

The problem with EAG: It's large (800-1,500 attendees), infrequent (a few times per year), and not exclusively top talent. Signal gets diluted. A first-time attendee can feel overwhelmed. Experienced attendees fill calendars early, newcomers get left out.

This event takes EAG's proven core mechanic and makes it:

- **Higher signal** (top 1% talent only, curated both sides)
- **More exclusive** (50 people, invite-only)
- **Monthly** (recurring cadence builds relationships and tracking over time)
- **With dedicated follow-up** (assigned champion for each participant, not just a survey)

Why BlueDot

BlueDot is the only organization with both assets needed to make this work:

1. **A curated talent database.** 7,000+ alumni with granular performance data. BlueDot doesn't just know who applied to courses. They know who performed best, who produced the strongest projects, who facilitators flagged as exceptional. No other organization in AI safety has this depth of talent signal.

2. **Institutional relationships on the demand side.** BlueDot already works directly with Coefficient Giving, MATS, Apart Research, Constellation, and hiring managers at Anthropic, Redwood, METR, and other safety orgs. They can get the decision-makers in the room.

The combination is the moat. Lots of organizations know talented people. Lots of organizations know institutional leaders. BlueDot is the only one that can curate both sides of the room with data-backed confidence, monthly, in SF.

The Thesis

Create the highest-signal talent marketplace in AI safety. Every attendee leaves with at least one concrete opportunity they wouldn't have found otherwise, whether that's a job lead, a funded project, a co-founder, or a program placement.

The curation is the product. The promise to both sides:

- **To talent:** "We'll connect you with the most influential people in safety so you can accelerate your career."
- **To institutions:** "We've done the screening. These are the people you should know about but don't."

---

2. Event Design: One Edition

Overview

- **Format:** Monthly in-person event, San Francisco
- **Duration:** 3 hours
- **Capacity:** ~50 attendees
- **Venue:** Mox (1680 Mission St, SF) - BlueDot already hosts SF Community Socials here. ~20k sq ft, AI safety coworking space.

Attendee Composition

Side A: Institutional Leaders (10-15 people) Representatives from the demand side of AI safety. Rotating each month to keep the event fresh.

Category	Example Orgs	What They Bring
Funders	Coefficient Giving, SFF, Open Phil	Funding signals, grant guidance
Safety orgs	Anthropic, Redwood, METR, DeepMind	Hiring leads, research direction
Programs	MATS, Constellation, Apart, CAIS	Program placements, mentorship
BlueDot leadership	Dewi, Li-Lian, team	Host, matchmaker, connector

Side B: Top-Tier Talent (30-35 people) Curated for talent quality and readiness. Sources:

Source	Selection Signal
BlueDot alumni (7,000+ pool)	Course performance, project quality, facilitator recommendations
MATS alumni	Research output, mentor evaluations
ATK network	Track record in their programs
Referrals from Side A	"Who should be in this room that isn't?"
Incubator Week applicants (not selected)	Still talented, just not ready for that format

Application and Curation

Invite-only, then apply. This is not an open event. BlueDot identifies the most promising people from their talent database and directly invites them to apply. The invite itself is the value proposition: "We'll connect you with the most influential people in safety so you can accelerate your career."

The application step (even for invited people) serves two purposes: (1) it collects the data needed for pre-matching, and (2) it filters for seriousness, since not everyone who's talented is ready or motivated to act right now.

Application Form - Side B (Talent)

Field	Type	Purpose
Name	Text	Identity
Email	Email	Contact
Current role and organization	Text	Background signal
Brief bio (100 words max)	Textarea	For other attendees to browse
What's your background? Highlight your strongest track record.	Textarea	Talent quality filter

Field	Type	Purpose
What area of AI safety are you most interested in?	Multi-select: Technical safety / AI governance / Biosecurity / Operations / Entrepreneurship / Other	Matching input
What are you looking for right now?	Multi-select: Full-time role / Project funding / Co-founder / Mentorship / Collaboration / Exploring options	Matching input + outcome tracking
What could you offer to someone else at this event?	Textarea	Two-way value signal
If you have a project or startup idea (even half-formed), describe it in 2-3 sentences.	Textarea (optional)	For funder matching
How did you hear about this event?	Dropdown	Outreach channel tracking
LinkedIn or personal site	URL	Additional vetting signal

Application Form - Side A (Institutional Leaders)

Field	Type	Purpose
Name	Text	Identity
Organization and role	Text	Identity
Brief bio (100 words max)	Textarea	For attendees to browse
What are you looking for?	Multi-select: Hiring for roles / Evaluating funding proposals / Scouting for programs / Meeting potential collaborators / Other	Matching input
What specific roles, project types, or profiles are you most interested in meeting?	Textarea	Matching precision

Field	Type	Purpose
What can you offer attendees?	Textarea	Value signal for matching

Curation criteria:

- Side B: Invite-only from BlueDot's talent database. Filtered by course performance, project quality, and facilitator recommendations. Invitees then apply with the form above. Select for: track record + seriousness + clear "ask."
- Side A: Direct outreach from BlueDot leadership. Must have decision-making power or strong influence over hiring, funding, or program admissions. No "observer" seats.

The 3-Hour Format

Time	Duration	Activity	Purpose
0:00-0:10	10 min	Arrival + welcome	Quick orientation: how the evening works, where to find your matches, logistics
0:10-1:25	75 min	**Pre-matched 1-on-1s (5 rounds x 15 min)**	High-value conversations with people selected for complementary needs. 12 min conversation + 3 min transition.
1:25-1:35	10 min	Break + refresh	Snacks, drinks, informal chat. Check the attendee board for self-selected matches.
1:35-2:40	65 min	**Self-selected 1-on-1s (4-5 rounds x 15 min)**	Attendees choose their own conversations based on what they've learned. Grab

Time	Duration	Activity	Purpose
			anyone. Swapcard facilitates requests.
2:40-2:55	15 min	**Next Steps capture**	Every attendee fills out a short form: top 1-2 action items from tonight. "I will email X about Y by [date]." This is the measurement mechanism.
2:55-3:00	5 min	Close	Thank attendees, remind about follow-up survey, done.

Why this structure works:

- **~10 conversations per person.** Matches EAG's target of ~10 new connections per attendee, compressed from a weekend to 3 hours.
- **Pre-matched first, self-selected second.** Ensures everyone has quality conversations (especially newcomers), then opens up for serendipity.
- **15-minute slots.** Long enough for a real conversation about opportunities, short enough to keep energy high and allow volume. Follow-ups happen after the event for deeper discussions. Unlike multi-day EAG events where meeting fatigue is a documented problem, this 3-hour format is a single high-intensity session. The short duration keeps energy high and prevents burnout.
- **Next Steps capture is the critical mechanism.** Converts conversations into commitments. Feeds the outcome tracking system.

Why Monthly Works

BlueDot runs monthly course cohorts across 4 concurrent courses (AGI Strategy, Technical AI Safety, Frontier AI Governance, Biosecurity). Each cohort produces participants rated by facilitators. The top-rated graduates from each cohort form the attendee pipeline for Side B. This means the event draws from a continuously replenishing pool each month rather than recycling the same people. Apart Research's monthly hackathon sprints demonstrate that consistent monthly cadence builds community habit and grows participation over time (from 118 signups to 766 over 18 months).

Pre-Event Matching Process

1. **Week before:** All attendee profiles visible on Swapcard. Attendees can browse and request meetings.
2. **3 days before:** BlueDot team reviews profiles and creates 5 pre-matched 1-on-1s per person based on complementary "looking for" / "can offer" fields. Prioritize talent-institution matches, but also include talent-talent matches for collaboration potential.
3. **Day of:** Attendees arrive with their schedule of 5 pre-matched meetings. Remaining 4-5 slots are self-selected during the event via Swapcard.

---

3. Metrics

North Star

Opportunities created: Job interviews, offers, hires, funded projects, and collaborations that trace back to connections made at this event.

Why Surveys Alone Aren't Enough

A case-control study of EAGxAustralia (Fodor & Tidmarsh, 2024) tested 20 measures at baseline and 6 months post-event, comparing attendees vs. non-attendees. They found no statistically significant differences. Standard post-event surveys may not capture real impact. This is why we track outcomes through dedicated follow-up (facilitators/champions), not just forms.

Three-Tier Metrics Framework

Tier 1: Attribution-tracked outcomes (the metrics that matter most) Tracked via dedicated facilitator follow-up calls + outcome check-in forms. These are the real measures of impact.

Metric	Tracked By	Target (initial hypothesis)
Job interviews traced to event connections	Facilitator calls + check-in form	3-5 per event
Job offers / hires traced to event connections	Facilitator calls + check-in form	1-3 per event

Metric	Tracked By	Target (initial hypothesis)
Funding applications submitted (including Rapid Grants)	Facilitator calls + check-in form	3-5 per event
Funding received	Facilitator calls + check-in form	1-2 per event
Collaborations or projects started	Facilitator calls + check-in form	2-4 per event
Program placements (MATS, etc.)	Facilitator calls + check-in form	1-2 per event

Note: These targets are initial hypotheses based on EAG connection-to-outcome ratios and the higher curation level of this event. The first 2-3 editions will calibrate these baselines. The important thing is having the tracking infrastructure in place to measure, not hitting specific numbers from day one.

Tier 2: Behavioral signals (post-event survey) Captured same evening. These predict whether Tier 1 outcomes will follow.

Metric	Source
"Plan changes" checklist: Did you resolve to apply for a role / submit a funding application / start a project / change career direction?	Post-event form
Counterfactual: "Would these plans have happened without this event?"	Post-event form
Number of follow-up meetings already scheduled	Post-event form
Impactful connections (defined as: "someone you'd feel comfortable reaching out to ask for a favor")	Post-event form
"Was this event more than 10x better than your counterfactual use of 3 hours?"	Post-event form

Tier 3: Event health (post-event survey) These tell us if the event itself is working, independent of outcomes.

Metric	Source	Target
Likelihood to Recommend (1-10)	Post-event form	>8
Matching quality (1-5)	Post-event form	>3.5
Would attend again	Post-event form	>80% yes
Meaningful connections per person	Post-event form	5+

Post-Event Feedback Form (sent same evening)

Field	Type	Purpose
How likely are you to recommend this event?	1-10 scale	Likelihood to Recommend (primary satisfaction metric, same as EAG)
How many new connections did you make that you'd feel comfortable reaching out to for a favor?	Number	EAG's definition of meaningful connections
Who was your most valuable conversation with and why?	Textarea	Identifies high-value attendees for future events
What specific opportunities came out of tonight?	Checkbox: Job lead / Funding lead / Co-founder match / Mentorship / Collaboration / Program referral / Rapid Grant idea / None yet	Outcome categorization
Did you make any plan changes as a result of tonight?	Checkbox: Apply for a role / Submit funding application / Start a project / Change career direction / Found an organization / None	EAG-style plan changes metric

Field	Type	Purpose
Would these plans have happened without this event?	Yes / Probably / Probably not / Definitely not	Counterfactual attribution
How many follow-up meetings have you already scheduled?	Number	Leading indicator of outcomes
Was this event more than 10x better than your counterfactual use of 3 hours?	Yes / No	High bar satisfaction check (from EAG)
Was the matching quality good?	1-5 scale	Matching algorithm feedback
Would you attend next month?	Yes / No / Maybe	Retention signal
Who should we invite next time that wasn't here?	Textarea	Referral pipeline
What would you change about the format?	Textarea	Iteration input

30/60/90 Day Follow-Up Check-In

Field	Type	Purpose
Since attending [Event Date], have any of the following happened as a result of connections made?	Checkbox: Applied for a role / Had a job interview / Received a job offer / Started a new role / Submitted a funding application / Received funding / Applied for Rapid Grant / Received Rapid Grant / Started a project or collaboration / Joined a program (MATS, etc.) / Found a co-founder / Ongoing mentorship / None yet	Outcome tracking

Field	Type	Purpose
If yes, briefly describe what happened.	Textarea	Qualitative data + testimonials
Which connection from the event led to this?	Text	Attribution tracking
Is there anything BlueDot could do to help you follow through?	Textarea	Ongoing support signal

Budget Estimate

Item	Monthly Cost	Notes
Venue (Mox)	$0-500	BlueDot already hosts at Mox; likely free or subsidized
Food/drinks	$500-750	Light catering for 50 people
Swapcard	~$50	$610/year Starter plan, amortized monthly
Printing/materials	$50-100	Name badges, schedules, signage
Staff time	Included	Part of Community Lead + team responsibilities
Total	**~$600-1,400/month**	**~$7,200-16,800/year**

At ~~$12-28 per attendee per event, this is far cheaper than EAG (~~$1,200/attendee) or even EA North (GBP 50/attendee). Unlike multi-day conferences, the cost scales linearly and predictably.

---

4. Outreach Strategy

Sourcing Side A (Institutional Leaders)

Channel	Approach
BlueDot leadership direct outreach	Dewi/Li-Lian personally invite key institutional contacts. High-touch, relationship-based.
Existing partnerships	Leverage co-hosting relationships with Apart, MATS, Constellation
SF AI safety ecosystem	Mox community, Bay Area AI Safety Meetup organizers
Rotating invitations	Different institutional reps each month to keep the event fresh and expand the network

Sourcing Side B (Top Talent)

Channel	Approach
BlueDot alumni database	Filter 7,000+ alumni by course performance, project quality, facilitator recommendations. Direct invite to top performers.
MATS alumni network	Partner with MATS to identify alumni who are job-seeking or exploring new directions
ATK / other talent pools	Direct partnership for cross-referrals
Referrals from Side A	Each institutional leader invited to refer 2-3 people. "Who should be in this room?"
BlueDot Slack community	Targeted announcements in relevant channels
Luma event page	Public-facing, but application-gated. Quality filter at application stage.

Outreach Cadence (Monthly Cycle)

Week	Activity
Week 1 (post-event)	Send follow-up surveys. Analyze feedback. Identify improvements.
Week 2	Confirm next month's Side A roster. Begin Side B sourcing. Open applications.
Week 3	Close applications. Curate attendee list. Set up Swapcard event.
Week 4	Pre-matching. Attendees browse profiles and schedule. Event day.

Part 2: Systems Design

System Overview

The system has five layers, designed to make the event repeatable and improve month over month.

Note: While Swapcard is simpler tools could handle 50-person logistics, its AI matchmaking and profile system reduce manual curation effort and scale cleanly if the event grows. At $610/year, the cost is negligible. Implementation should also build on BlueDot's existing infrastructure where possible. The first step after onboarding would be auditing current internal tools (e.g., the BlueDot monorepo, Airtable bases, Slack integrations) and adapting this design to leverage what already exists.

Layer 1: Event Operations (Swapcard + Luma)

- **Swapcard:** Attendee profiles, 1-on-1 scheduling, AI matchmaking, in-event logistics
- **Luma:** Event page, RSVP, reminders, basic CRM

Layer 2: Data Collection (Forms)

- Application form (pre-event, via Swapcard)
- Next Steps form (filled by participants at event close)
- Post-event feedback form (sent same evening)

- 30/60/90 day outcome check-ins (separate forms, sent automatically)

Layer 3: Individual Participant Tracking (THE KEY DIFFERENTIATOR)

This is what makes the event actually work. The event creates connections. The follow-up converts connections into outcomes.

Every participant gets:

- **An assigned facilitator** (responsible for their follow-up)
- **Monthly check-in calls** (biweekly for high-priority participants with active opportunities)
- **A running profile** tracking their goal, blockers, actions taken, and outcome status

The facilitator's job on each call:

1. What happened since last time? (track progress)
2. What's blocking you? (identify where BlueDot can help)
3. Who should you talk to next? (make warm introductions)
4. What's your next concrete step? (maintain momentum)

Why this matters: If these are truly the top talent in AI safety, they deserve dedicated follow-up. A 30-minute call every month is a small investment. The return is a placement, a funded project, or a collaboration that wouldn't have happened otherwise. This is the mechanism that turns a monthly event into a talent acceleration program.

Staffing model: BlueDot already has a facilitator network (alumni who facilitate course cohorts, 0.5-1 FTE each). These facilitators can extend their role to track event participants. This is a natural fit: they already have relationship-building skills and know the BlueDot ecosystem. Additionally, partnerships with organizations like 80,000 Hours or SuccessIf (who do similar career tracking) could share the load.

With ~35 Side B participants per event and some returning month over month, the active tracking pool might be 50-80 people at steady state. At ~1 call/month per person, that's 50-80 calls/month, roughly 25-40 hours. Distributed across 3-5 facilitators, this is ~8 hours/week per person, a manageable part-time commitment.

Tracked in Airtable: The Participant Tracking table records each person's assigned champion, call cadence, current goal, key connections, blockers, actions taken, outcome status, and priority level.

Layer 4: Outcome Tracking (Airtable)

Central database connecting:

- Next steps committed to at events
- Outcomes reported at 30/60/90 days via check-in forms
- Champion call notes and progress updates
- Feedback scores and qualitative input
- Longitudinal tracking across multiple events

Layer 5: The Improvement Engine (How This Gets Better Every Month)

The assignment asks for a system that helps "improve this event every month." This section makes that explicit.

The core principle: Every event generates data. That data answers specific questions. Those answers drive specific changes. We measure whether the changes worked.

Six Improvement Dimensions

1. Matching Quality: "Are we pairing the right people?"

- **Data source:** Matching Quality score (1-5) from Post-Event Feedback + "Most Valuable Conversation" responses
- **Question to ask:** Do pre-matched 1-on-1s score higher than self-selected ones? Which match criteria (same interest area? complementary roles? same career stage?) predict the best conversations?
- **Action:** Adjust matching weights in Swapcard. If "same interest area" matches score poorly but "complementary role" matches score well, shift the algorithm.
- **Success signal:** Avg matching quality trending upward month over month.

2. Attendee Composition: "Do we have the right people in the room?"

- **Data source:** "Most Valuable Conversation" (who gets named repeatedly?), Outcome Tracking (which attendee types produce the most outcomes?), "Who Should We Invite Next Time?"
- **Question to ask:** Are we getting enough institutional leaders? Are certain types of talent (e.g., technical vs. policy) over/underrepresented relative to demand? Are returning attendees still getting value, or does the room need fresh faces?
- **Action:** Adjust the invite list. If hiring managers from safety orgs are consistently the most valuable Side A attendees, get more of them. If attendees from a certain source (e.g., MATS) consistently produce outcomes, increase that pipeline.
- **Success signal:** Outcome rate per event trending upward. "Would attend again" staying above 80%.

3. Format and Timing: "Is the structure working?"

- **Data source:** "What Would You Change?" from Post-Event Feedback + observation by the host during the event
- **Question to ask:** Are 15-min slots too short or too long? Is the break well-placed? Do people use all self-selected slots? Is the Next Steps capture rushed?
- **Action:** Test changes. If multiple people ask for longer conversations, try 20 min x 8 rounds next month. If self-selected slots are underused, add more pre-matched rounds. Run A/B tests across months.
- **Success signal:** NPS trending upward. Connections per person stable or increasing.

4. Follow-Up Effectiveness: "Are champion calls converting connections into outcomes?"

- **Data source:** Participant Tracking table (outcome status, actions taken, blockers), Outcome Tracking (30/60/90 day results)
- **Question to ask:** What percentage of participants with an active champion convert to a positive outcome within 90 days? What are the most common blockers? Are champions making enough warm introductions?
- **Action:** If outcome rate is low despite good event feedback, the problem is follow-up, not the event. Train champions on more effective calls. Increase call frequency for high-priority participants. Create a playbook for common blockers.
- **Success signal:** 90-day outcome rate increasing. Blocker patterns becoming less frequent as playbooks address them.

5. Outreach and Curation: "Are we finding the best people?"

- **Data source:** "How did you hear about this?" from applications, post-event quality ratings correlated with source
- **Question to ask:** Which sourcing channels produce the highest-quality attendees? Are direct invites from the BlueDot database outperforming referrals, or vice versa? What's the acceptance rate on invitations?
- **Action:** Double down on the channels that produce the best attendees. If referrals from Side A ("who should be in this room?") consistently produce the best Side B attendees, systematize that. If a talent pool isn't converting, drop it.
- **Success signal:** Application quality increasing. Less curation effort needed to fill 50 high-quality seats.

6. Institutional Value: "Are Side A leaders getting enough value to keep coming?"

- **Data source:** Side A-specific feedback (can add 2-3 questions to their version of the post-event form), repeat attendance rates for institutional leaders
- **Question to ask:** Are they meeting people they want to hire or fund? Is the event worth 3 hours of their time monthly? What would make it more valuable for them?

- **Action:** If Side A retention drops, diagnose why. They might want higher-caliber talent, more preparation on who they'll meet, or a different format (e.g., 2 hours not 3). Their continued attendance is critical, so prioritize their feedback.
- **Success signal:** Side A repeat attendance rate above 60%. New institutional leaders joining through word-of-mouth.

The Monthly Retrospective Process

After each event, the host produces a **one-page Event Retrospective** (stored in the Events table "Retro Notes" field):

1. **Review quantitative data** from the Proxy Metrics Dashboard. Compare to previous months. Flag any metric that moved significantly (positive or negative).
2. **Read all qualitative feedback.** Identify the top 3 themes from "What Would You Change?" and "Most Valuable Conversation."
3. **Check champion call notes.** What patterns are emerging in participant tracking? What are the most common blockers?
4. **Set one specific improvement action for next month.** Not three. One. Test it, measure it, then move to the next thing. Small, deliberate changes compound over time.
5. **Update the Proxy Metrics Dashboard** with the "Top Improvement Action" field so there's a running record of what was changed and when.

What This Looks Like Over 6 Months

Month	Likely Focus	Example Improvement
1	Logistics	Get the basics right: venue, timing, form flow, Swapcard setup
2	Matching	Adjust matching criteria based on Month 1 feedback on conversation quality
3	Follow-up	Refine champion call process based on what's working/not in first 60 days of tracking
4	Composition	Rebalance Side A / Side B ratio or profiles based on outcome data
5	Outreach	Scale up best sourcing channels, drop underperforming ones

Month	Likely Focus	Example Improvement
6	Format	Test a structural change (timing, group activities, theme nights) based on accumulated feedback

Over time, this creates a dataset that answers the strategic question: "What predicts whether a connection made at our event turns into a real opportunity?" That prediction model, refined month after month, becomes BlueDot's most valuable asset for talent acceleration. It's not just an event. It's a learning machine.

Airtable Base: How to Use It

Link: [BlueDot Talent Marketplace Airtable Base]

The Airtable base has 6 tables. Here's the workflow:

Before Each Event

1. **Events table:** Create a new row for the upcoming edition (name, date, venue, edition number)

At the Event (last 15 minutes)

2. **Next Steps table:** Share the Airtable form link. Each participant fills in their top 1-2 action items with target dates. Takes 3-5 minutes.

Same Evening (within 2 hours of event end)

3. **Post-Event Feedback table:** Send the feedback form link via email (automate through Luma or Airtable). Captures NPS, connection count, matching quality, opportunities emerged, format suggestions, and referrals.

Week After the Event

4. **Participant Tracking table:** Create or update a row for each participant. Assign a BlueDot champion. Set call cadence and next call date. Record their current goal and priority level.
5. **Events table:** Fill in the aggregated metrics (NPS, connections, attendee counts) and write the retro notes.

6. **Proxy Metrics Dashboard:** Summarize the event's performance in one row. Compare to previous months.

30/60/90 Days After

7. **Outcome Tracking table:** Send the check-in form at each interval. Record outcomes reported. Update Participant Tracking with any status changes.
8. **Proxy Metrics Dashboard:** Update the lagging indicators (job apps, interviews, offers, funding, collaborations) as outcome data comes in.

Monthly Improvement Review

9. **Proxy Metrics Dashboard:** Compare the current event to all previous events. Identify trends. Set the "Top Improvement Action" for next month.

Assumptions and Uncertainties

Assumption	**Confidence**	**If Wrong**
Mox is available monthly for this format	High (BlueDot already hosts there)	Alternative: BlueDot's own SF office once open
10-15 institutional leaders will attend monthly	Medium	May need stronger incentives or rotating smaller groups
Swapcard works at 50-person scale	Medium (designed for larger events)	Fallback: Airtable forms + manual matching
15-minute 1-on-1s are long enough for meaningful conversations	Medium	Could adjust to 20 min with fewer rounds
Top talent will self-select into applications	Medium	May need more active recruiting from the database
30/60/90 day follow-up response rates will be high enough for tracking	Low-Medium	Incentivize with: continued event access contingent on feedback

---

What Makes This Different

Existing Format	Gap This Event Fills
EAG	Too large, not exclusive enough. Signal diluted across 800+ people of varying quality levels.
BlueDot Incubator Week	Too intensive (5 days), too selective (7 people), too rare (on hiatus). Not a pipeline, a one-off.
Apart Research hackathons	Output-focused (research papers). Good for demonstrating capability, not for matching to opportunities.
Mox socials	Unstructured. No curation, no matching, no outcome tracking.
Bay Area AI Safety Meetup	Speaker-focused, not matching-focused.

This event is the only monthly, curated, outcome-tracked talent marketplace for AI safety in SF.

Participant Tracker Guide

Why Individual Tracking Matters

The event creates connections. Tracking converts connections into outcomes. Without dedicated follow-up, even the best-matched 1-on-1 fades into a LinkedIn connection that never goes anywhere.

The research backs this up: a case-control study of EAGxAustralia found no statistically significant impact differences between attendees and non-attendees at 6 months using standard surveys. Surveys alone don't capture real impact. Direct, relationship-based follow-up does.

Who Does the Tracking

Facilitators. BlueDot already has a facilitator network: alumni who facilitate course cohorts (0.5-1 FTE each, 2 weeks of training). These people already have relationship-building skills and know the BlueDot ecosystem. Extending their role to track event participants is a natural fit.

Each facilitator is assigned 10-15 participants. At one 30-minute call per participant per month, that's 5-8 hours/month per facilitator, a manageable part-time commitment.

Potential partnerships: 80,000 Hours does similar career tracking. SuccessIf also operates in this space. Joint efforts could share the load and avoid duplicating outreach to the same people.

The Tracking Process

After Each Event (Week 1)

1. **Create or update Participant Tracking records** in Airtable for every Side B attendee
2. **Assign a facilitator** to each new participant. Try to keep the same facilitator across events for continuity.
3. **Set the participant's current goal** based on their application form ("looking for" field) and Next Steps form
4. **Set priority level:**
    - **High** (biweekly calls): Has a specific opportunity in motion (interview scheduled, grant application started, co-founder identified)
    - **Medium** (monthly calls): Has a direction but no specific opportunity yet
    - **Low** (as needed): Exploring, not ready to act
5. **Schedule the first call** within 1-2 weeks of the event, while conversations are fresh

Monthly Check-In Call (30 minutes)

Opening (5 min): "Since we last spoke, what's happened? Any conversations from the event that went somewhere?"

Progress review (10 min):

- Update "Actions Taken" in Airtable with anything they've done
- Update "Outcome Status" if it's changed
- Check: did they follow through on the Next Steps they wrote down at the event?

Blocker identification (5 min):

- "What's preventing you from taking the next step?"

- Common blockers and how to help:
    - "I don't know who to talk to" -> Make a warm introduction from BlueDot's network
    - "I'm not sure if my idea is good enough to apply for funding" -> Connect them with a past Rapid Grant recipient or offer to review their draft
    - "I'm waiting to hear back from X" -> Follow up on their behalf if appropriate
    - "I'm not sure which path to take (role vs. project vs. startup)" -> Suggest they attend next month's event with a more focused "looking for"

Next steps (5 min):

- Set 1-2 concrete action items with dates
- Update "Current Goal" and "Next Call Date" in Airtable

Warm introductions (5 min):

- Based on what you've heard, is there someone in BlueDot's network who could help?
- Log any intros made in "Key Connections Made"

Outcome Recording

When a participant reports a concrete outcome (got an interview, received funding, started a project):

1. Update **Outcome Status** in Participant Tracking (e.g., "Active - in process" -> "Placed (role)")
2. Create a record in the **Outcome Tracking** table with details
3. Ask: "Can we share this as a testimonial?" (builds the case for future events)
4. Note which event connection led to this outcome (attribution data)

When to Adjust Priority

Signal	**Action**
Participant gets an interview or funding lead	Upgrade to High, switch to biweekly calls
Participant hasn't responded to 2 check-ins	Downgrade to Low, switch to async (Slack DM)
Participant achieves their goal (hired, funded, etc.)	Mark as completed, move to quarterly check-ins
Participant says they're no longer pursuing AI safety	Mark as Inactive, stop calls, note reason

Facilitator Dashboard (Airtable View)

Create a filtered view in the Participant Tracking table for each facilitator showing only their assigned participants, sorted by Next Call Date. This is their weekly to-do list.

Key fields to surface in the view:

- Participant Name
- Priority (color-coded)
- Current Goal
- Next Call Date
- Last Call Date
- Outcome Status
- Blockers

Monthly Facilitator Sync (30 min, all facilitators)

Once a month, facilitators meet to:

1. Share patterns ("three of my participants are stuck on the same blocker")
2. Cross-refer participants to each other's networks
3. Identify systemic issues to feed back into the event improvement engine
4. Celebrate wins (placements, funding, collaborations)

This sync is where individual tracking becomes organizational intelligence.

Appendix A: Methodology (How I Built This Submission)

I want to be transparent about the process I followed. Here's the step-by-step approach:

Step 1: Deep Research on BlueDot (before the clock started)

Before starting the timed work test, I compiled a comprehensive research document on BlueDot Impact: their theory of change, course portfolio, community infrastructure, events program, career services, team, funding, and the SF AI safety ecosystem. This gave me a solid foundation of what BlueDot already does, what they value, and where the gaps are. (See Appendix B.)

Step 2: Identify the Bottleneck (first ~45 minutes)

Rather than jumping to event format, I started with: "What specific problem should this event solve?" I mapped the AI safety talent pipeline into stages (awareness, motivation, skills, first role, connection, founder pipeline) and evaluated where a monthly 50-person event would have the most leverage. This was a deliberate brainstorming process, testing and discarding several framings before landing on the right one.

Key decision point: I considered several event concepts:

- A "founder meets funder" pitch night (too close to Incubator Week)
- A "problem discovery engine" for builders (too research-oriented)
- A career matching event like HackerX (too transactional)
- A curated talent marketplace (the winner)

The curated talent marketplace won because it addresses the actual bottleneck (talent exists but isn't meeting opportunities efficiently), plays to BlueDot's unique assets (talent database + institutional relationships), and has a proven model to build on (EAG's 1-on-1 matching).

Step 3: Research Existing Models (~30 minutes)

I researched six areas to inform the design:

1. **EAG events** - How their 1-on-1 system works, metrics, retrospectives, failure modes
2. **Apart Research sprints** - Monthly cadence model, scaling trajectory, co-hosting approach
3. **ControlConf** - Conference format, workshop model
4. **Swapcard** - Platform features, how EAG uses it, alternatives for smaller events
5. **AI safety talent pipeline bottlenecks** - Academic studies, hiring manager interviews, fellowship outcome data
6. **Structured career-matching formats** - YC Demo Day, HackerX, reverse job fairs, Recurse Center model

(See Appendix C for the full research document.)

Step 4: Design the Event (~30 minutes)

With the bottleneck identified and models researched, I designed the event format, attendee composition, application process, and 3-hour schedule. Key design decisions were informed by evidence:

- 15-minute 1-on-1 slots (balancing EAG's 25-minute standard with the 3-hour constraint)
- Pre-matched + self-selected rounds (from EAG's newcomer problem research)

- Invite-only curation (from HackerX's pre-vetting model)
- Next Steps capture (from the finding that surveys alone don't predict real outcomes)

Step 5: Design the System (~45 minutes)

Built the Airtable base with 6 tables and designed the 5-layer system. The key insight was that the improvement engine needs to be explicit, not assumed: every dimension of the event (matching, composition, format, follow-up, outreach, institutional value) has its own data source, question, action, and success signal.

The participant tracking system (facilitators doing monthly calls) emerged from the realization that the event alone doesn't produce outcomes, but dedicated follow-up does.

Step 6: Red Team and Iterate (~30 minutes)

Ran an adversarial review of the full submission, identifying 15 findings across operational feasibility, metric validity, differentiation, and presentation. Applied fixes for the high-priority items (champion system feasibility, budget, metrics as hypotheses, monthly cadence justification, unique depth demonstration).

Tools Used

- **Claude Code** for brainstorming, research, drafting, and iteration
- **Exa search** for web research on EAG events, Apart Research, talent pipeline studies
- **Airtable API** to build the base programmatically
- **Google Docs** for the final submission

Appendix B: BlueDot Impact Research

A comprehensive research document compiled before the work test. Covers: organization overview, theory of impact, core beliefs, history, course portfolio, community infrastructure, events program, career services, incubator programs, team and funding, the Community Lead role, the SF AI safety ecosystem, and gaps/opportunities.

Full document available at: [link to bluedot-research.md if hosted, or included below]

Key findings that shaped this submission:

- **Workforce gap:** ~2,000 people currently work full-time on AI safety. BlueDot's target: 100,000 by 2030.

- **Course impact:** 5% of participants worked on AI safety pre-course, 37% afterward (2022 Alignment Course data).
- **Hiring needs (Li-Lian Ang, Feb 2026):** Orgs need people who can "own and lead work independently" and "care deeply about AI risk." Strong track records matter more than AI safety experience.
- **Coefficient Giving (March 2026):** More funding appetite than good applications. Capacity-building is undervalued.
- **SF ecosystem gap:** No dedicated AI safety talent-matching or career event exists in SF.
- **BlueDot's event philosophy:** Structured, not just social. Career-oriented. Data-driven. Community-led hosting model.
- **EA North proof of concept:** 35 people, one day, GBP 50/person, 4.56 connections per person, 12x more cost-effective than flagship EAG.

Appendix C: Event Design Research

A research document covering six areas that informed the event design. Compiled during the work test using web research across EA Forum, LessWrong, Apart Research, CEA publications, and other sources.

Key findings by topic:

Apart Research Sprint Model

- Monthly cadence builds habit: grew from 118 to 766 signups over 18 months
- Co-hosting multiplies reach (Redwood, BlueDot, METR, Goodfire)
- 22 peer-reviewed publications from hackathon outputs
- The hackathon-to-fellowship pipeline is the retention mechanism

EAG Events and 1-on-1 Scheduling

- ~10 new connections per attendee at EAGx events
- Cost per connection: ~$70 (EAGx) to ~$130 (flagship EAG)
- Common failure modes: meeting exhaustion (10-15/day), transactional feel, newcomers getting left out
- 1-on-1s are the core value, not talks. CEA explicitly coaches attendees on this.

Swapcard

- AI matchmaking, profile-based scheduling, analytics
- $610/year Starter plan covers 200 attendees/year

- Overkill feature-wise for 50 people, but the matchmaking engine saves manual curation effort

AI Safety Talent Pipeline Bottlenecks

- Primary bottleneck: training to first role (Stage 2)
- 10-30% of fellowship alumni cycle into another fellowship (Clay dataset, 600+ alumni)
- Senior talent scarcity creates a cascade: without enough senior people, juniors can't be mentored
- Automation Paradox: AI tools raise the bar for human contribution rather than lowering it (Teichman study)
- 80K Hours survey: "attend events and get to know people" is top career advice from AI safety org leaders

Career-Matching Event Formats

- HackerX: pre-vetting both sides is the key to high signal-to-noise
- YC Demo Day: extreme curation + time compression + pre-existing trust
- Reverse job fairs: putting candidates in the active position
- Fonzi AI weekly match days: recurring cadence keeps both sides engaged
- EA North: proof that 25-75 person regional events work at 12x better cost-effectiveness than EAG